

Title: Linear Regression vs Polynomial Regression

Goal 🎯: Find which model works best without overfitting using Advertising dataset.



Lets start with importing the libraries and loading the dataset

```
In [21]: import numpy as np  
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
In [22]: advertising_df = pd.read_csv('Advertising.csv')  
df = advertising_df.copy()  
df.head()
```

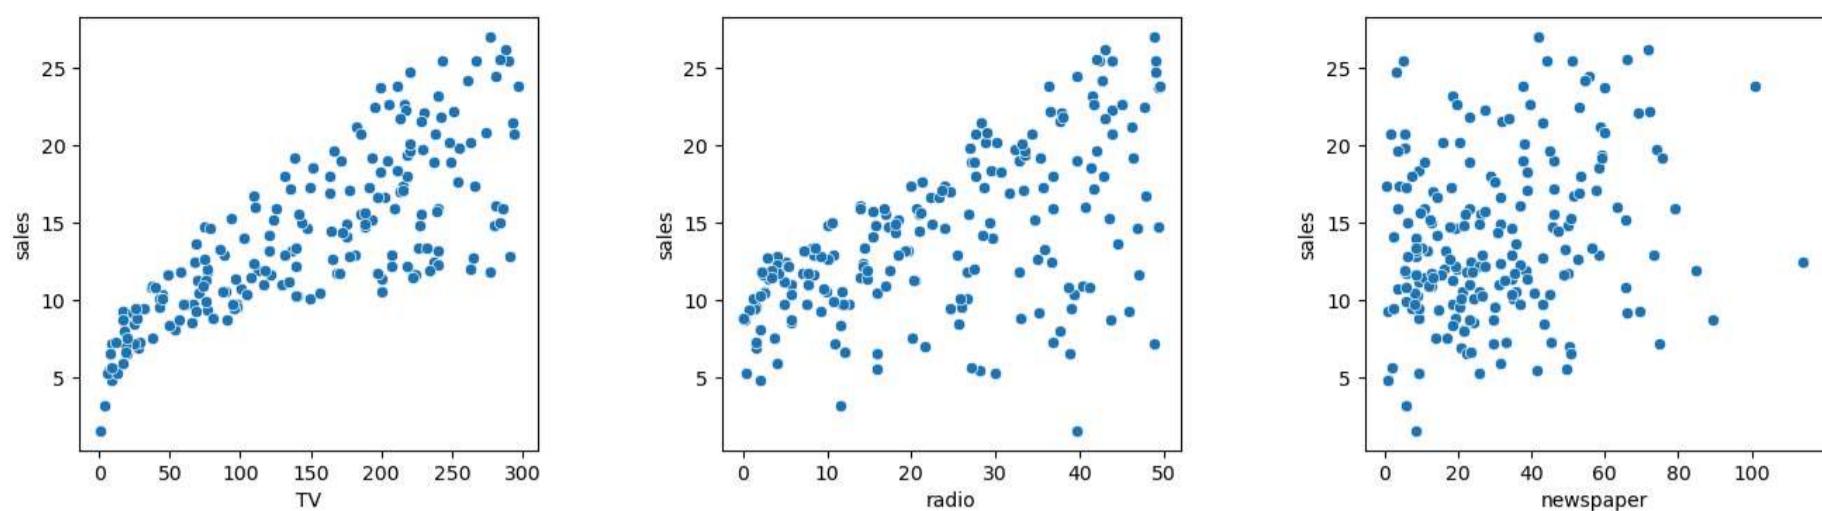
```
Out[22]:   TV  radio  newspaper  sales  
0    230.1    37.8      69.2    22.1  
1     44.5    39.3      45.1    10.4  
2     17.2    45.9      69.3     9.3  
3    151.5    41.3      58.5    18.5  
4    180.8    10.8      58.4    12.9
```

Lets check if the data is clean before we start the comparison.

```
In [23]: df.info()  
  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 200 entries, 0 to 199  
Data columns (total 4 columns):  
 #   Column      Non-Null Count  Dtype     
---  --          --          --  
 0   TV          200 non-null    float64  
 1   radio        200 non-null    float64  
 2   newspaper    200 non-null    float64  
 3   sales        200 non-null    float64  
 dtypes: float64(4)  
 memory usage: 6.4 KB
```

```
In [24]: df.isna().sum()  
  
Out[24]: TV      0  
radio     0  
newspaper 0  
sales     0  
dtype: int64
```

```
In [25]: plt.figure(figsize=(16,4))  
plt.subplot(1,3,1)  
sns.scatterplot(data=df,x='TV',y='sales')  
  
plt.subplot(1,3,2)  
sns.scatterplot(data=df,x='radio',y='sales')  
  
plt.subplot(1,3,3)  
sns.scatterplot(data=df,x='newspaper',y='sales')  
  
plt.subplots_adjust(wspace=0.4, hspace=0.4)
```



```
In [26]: df.describe().T
```

| | count | mean | std | min | 25% | 50% | 75% | max |
|------------------|-------|----------|-----------|-----|--------|--------|---------|-------|
| TV | 200.0 | 147.0425 | 85.854236 | 0.7 | 74.375 | 149.75 | 218.825 | 296.4 |
| radio | 200.0 | 23.2640 | 14.846809 | 0.0 | 9.975 | 22.90 | 36.525 | 49.6 |
| newspaper | 200.0 | 30.5540 | 21.778621 | 0.3 | 12.750 | 25.75 | 45.100 | 114.0 |
| sales | 200.0 | 14.0225 | 5.217457 | 1.6 | 10.375 | 12.90 | 17.400 | 27.0 |

The data is cleaned where there is no null-values nor zero values that doesn't make any sense.



It is time to train the models.

```
In [27]: X = df.drop('sales',axis=1)
y = df['sales']
```

```
In [28]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=101)
```

First, lets start with Linear Regression.

```
In [29]: from sklearn.linear_model import LinearRegression
lr_model = LinearRegression()
lr_model.fit(X_train,y_train)
```

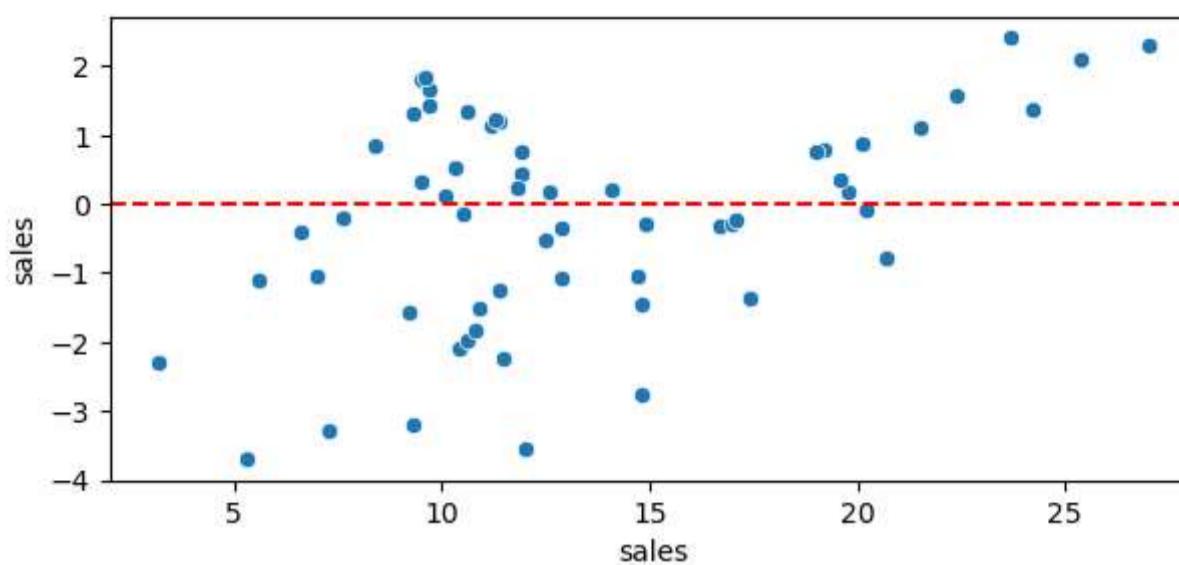
```
Out[29]: 
  ▾ LinearRegression ⓘ ?
```

```
In [30]: test_prediction = lr_model.predict(X_test)
```

I want to visualize using residual plot to make sure that Linear Regression can be used in this dataset.

```
In [31]: test_residuals = y_test - test_prediction
plt.figure(figsize=(7,3))
sns.scatterplot(x=y_test,y=test_residuals)
plt.axhline(y=0,color='red',ls='--')
```

```
Out[31]: <matplotlib.lines.Line2D at 0x255b9a9f4d0>
```



Now it is time for Polynomial Regression.

```
In [32]: from sklearn.preprocessing import PolynomialFeatures
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
```

I'm going to use for loop to check which degree is the best to use.

```
In [33]: from sklearn.metrics import r2_score, root_mean_squared_error
train_rmse_errors = []
test_rmse_errors = []
train_r_errors = []
test_r_errors = []

for i in range(1,10):
    poly_converter = PolynomialFeatures(degree=i, include_bias=False)
    poly_features = poly_converter.fit_transform(X)
    X_train, X_test, y_train, y_test = train_test_split(poly_features, y, test_size=0.3, random_state=101)
    X_train = scaler.fit_transform(X_train)
    X_test = scaler.transform(X_test)
    pr_model = LinearRegression()
    pr_model.fit(X_train, y_train)
    train_pred = pr_model.predict(X_train)
    test_pred = pr_model.predict(X_test)
    train_rmse = root_mean_squared_error(y_train, train_pred).round(2)
    test_rmse = root_mean_squared_error(y_test, test_pred).round(2)
    train_r = round(r2_score(y_train, train_pred), 2)
    test_r = round(r2_score(y_test, test_pred), 2)
    train_rmse_errors.append(train_rmse)
    test_rmse_errors.append(test_rmse)
    train_r_errors.append(train_r)
    test_r_errors.append(test_r)
    print(f'Train RMSE for degree {i}: {train_rmse}')
    print(f'Test RMSE for degree {i}: {test_rmse}')
    print(f'Train R^2 for degree {i}: {train_r}')
    print(f'Test R^2 for degree {i}: {test_r}')
    print()
```

Train RMSE for degree 1: 1.73

Test RMSE for degree 1: 1.52

Train R^2 for degree 1: 0.89

Test R^2 for degree 1: 0.92

Train RMSE for degree 2: 0.59

Test RMSE for degree 2: 0.66

Train R^2 for degree 2: 0.99

Test R^2 for degree 2: 0.98

Train RMSE for degree 3: 0.43

Test RMSE for degree 3: 0.58

Train R^2 for degree 3: 0.99

Test R^2 for degree 3: 0.99

Train RMSE for degree 4: 0.35

Test RMSE for degree 4: 0.51

Train R^2 for degree 4: 1.0

Test R^2 for degree 4: 0.99

Train RMSE for degree 5: 0.25

Test RMSE for degree 5: 2.58

Train R^2 for degree 5: 1.0

Test R^2 for degree 5: 0.76

Train RMSE for degree 6: 0.19

Test RMSE for degree 6: 7.23

Train R^2 for degree 6: 1.0

Test R^2 for degree 6: -0.85

Train RMSE for degree 7: 0.11

Test RMSE for degree 7: 113.15

Train R^2 for degree 7: 1.0

Test R^2 for degree 7: -452.45

Train RMSE for degree 8: 0.0

Test RMSE for degree 8: 1396.85

Train R^2 for degree 8: 1.0

Test R^2 for degree 8: -69110.83

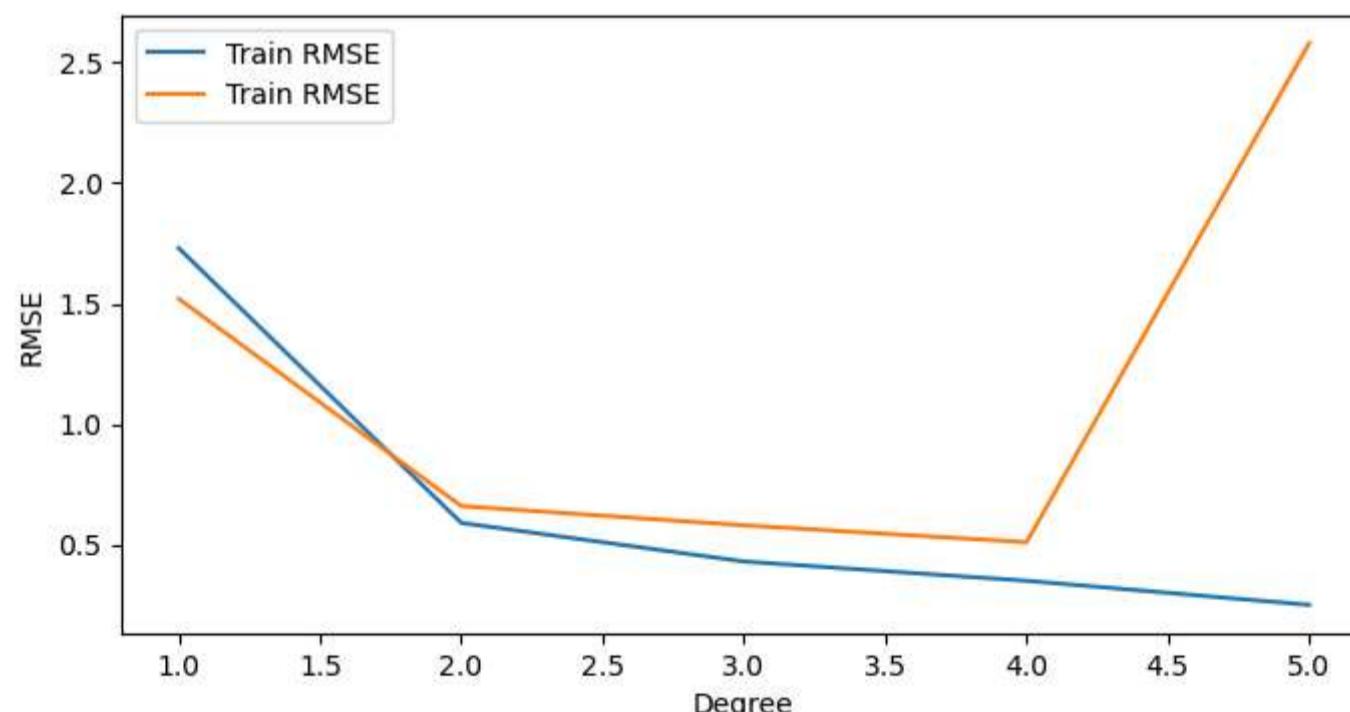
Train RMSE for degree 9: 0.0

Test RMSE for degree 9: 2342.54

Train R^2 for degree 9: 1.0

Test R^2 for degree 9: -194368.78

```
In [34]: plt.figure(figsize=(8,4))
plt.plot(range(1,6),train_rmse_errors[:5],label='Train RMSE')
plt.plot(range(1,6),test_rmse_errors[:5],label='Test RMSE')
plt.xlabel('Degree')
plt.ylabel('RMSE')
plt.legend()
plt.show()
```



From the results, I decided to use the 3rd degree since it gave us the best results while not being overfit.

```
In [35]: final_poly_converter = PolynomialFeatures(degree=3, include_bias=False)
poly_features = final_poly_converter.fit_transform(X)
X_train, X_test, y_train, y_test = train_test_split(poly_features, y, test_size=0.3, random_state=101)
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
final_pr_model = LinearRegression()
```

```
final_pr_model.fit(X_train,y_train)
final_pr_pred = final_pr_model.predict(X_test)
```

If we are still not sure if we have overfitting, lets use regularizations to prevent from overfitting. I will start with Ridge Regression with Cross Validation to give us the best alpha to use.

```
In [36]: from sklearn.linear_model import RidgeCV
# help(RidgeCV)
ridge_cv_model = RidgeCV(alphas=(0.1,1.0,10.0),cv=10,)
ridge_cv_model.fit(X_train,y_train)
ridge_cv_pred_values = ridge_cv_model.predict(X_test)
```

```
In [37]: poly_converter.get_feature_names_out()[0:19]
```

```
Out[37]: array(['TV', 'radio', 'newspaper', 'TV^2', 'TV radio', 'TV newspaper',
   'radio^2', 'radio newspaper', 'newspaper^2', 'TV^3', 'TV^2 radio',
   'TV^2 newspaper', 'TV radio^2', 'TV radio newspaper',
   'TV newspaper^2', 'radio^3', 'radio^2 newspaper',
   'radio newspaper^2', 'newspaper^3'], dtype=object)
```

```
In [38]: ridge_cv_model.coef_
```

```
Out[38]: array([ 5.40769392,  0.5885865 ,  0.40390395, -6.18263924,  4.59607939,
   -1.18789654, -1.15200458,  0.57837796, -0.1261586 ,  2.5569777 ,
   -1.38900471,  0.86059434,  0.72219553, -0.26129256,  0.17870787,
   0.44353612, -0.21362436, -0.04622473, -0.06441449])
```

The other regularization to use is the Lasso Regression with Cross Validation.

```
In [39]: from sklearn.linear_model import LassoCV
# help(LassoCV)
lasso_cv_model = LassoCV(cv=10,eps=0.001,n_alphas=100,max_iter=1000000)
lasso_cv_model.fit(X_train,y_train)
lasso_pred = lasso_cv_model.predict(X_test)
```

```
In [40]: poly_converter.get_feature_names_out()[0:19]
```

```
Out[40]: array(['TV', 'radio', 'newspaper', 'TV^2', 'TV radio', 'TV newspaper',
   'radio^2', 'radio newspaper', 'newspaper^2', 'TV^3', 'TV^2 radio',
   'TV^2 newspaper', 'TV radio^2', 'TV radio newspaper',
   'TV newspaper^2', 'radio^3', 'radio^2 newspaper',
   'radio newspaper^2', 'newspaper^3'], dtype=object)
```

```
In [41]: lasso_cv_model.coef_
```

```
Out[41]: array([ 4.86023329,  0.12544598,  0.20746872, -4.99250395,  4.38026519,
   -0.22977201, -0.        ,  0.07267717, -0.        ,  1.77780246,
   -0.69614918, -0.        ,  0.12044132, -0.        , -0.        ,
   -0.        ,  0.        ,  0.        , -0.        ])
```

Last one is Elastic-Net, which is mix between Ridge and Lasso Regressions.

```
In [42]: from sklearn.linear_model import ElasticNetCV
elastic_cv_model = ElasticNetCV(l1_ratio=[.1,.2,.3,.4,.5,.6,.7,.8,.9,.95,1],eps=0.001,cv=10,n_alphas=100,max_iter=1000000)
elastic_cv_model.fit(X_train,y_train)
elastic_pred = elastic_cv_model.predict(X_test)
```

```
In [43]: elastic_cv_model.l1_ratio_
```

```
Out[43]: 1.0
```

```
In [44]: poly_converter.get_feature_names_out()[0:19]
```

```
Out[44]: array(['TV', 'radio', 'newspaper', 'TV^2', 'TV radio', 'TV newspaper',
   'radio^2', 'radio newspaper', 'newspaper^2', 'TV^3', 'TV^2 radio',
   'TV^2 newspaper', 'TV radio^2', 'TV radio newspaper',
   'TV newspaper^2', 'radio^3', 'radio^2 newspaper',
   'radio newspaper^2', 'newspaper^3'], dtype=object)
```

```
In [45]: elastic_cv_model.coef_
```

```
Out[45]: array([ 4.86023329,  0.12544598,  0.20746872, -4.99250395,  4.38026519,
   -0.22977201, -0.        ,  0.07267717, -0.        ,  1.77780246,
   -0.69614918, -0.        ,  0.12044132, -0.        , -0.        ,
   -0.        ,  0.        ,  0.        , -0.        ])
```

Lastly, it is time to evaluate the models.

```
In [46]: print('Linear Regression:')
print(f'RMSE: {root_mean_squared_error(y_test,test_prediction).round(2)}')
print(f'R^2: {round(r2_score(y_test,test_prediction),2)}')
print()
print('Polynomial Regression, Degree 3:')
print(f'RMSE: {root_mean_squared_error(y_test,final_pr_pred).round(2)}')
print(f'R^2: {round(r2_score(y_test,final_pr_pred),2)}')
print()
print('Ridge CV Regression:')
print(f'RMSE: {root_mean_squared_error(y_test,ridge_cv_pred_values).round(2)}')
print(f'R^2: {round(r2_score(y_test,ridge_cv_pred_values),2)}')
print()
print('Lasso CV Regression:')
print(f'RMSE: {root_mean_squared_error(y_test,lasso_pred).round(2)}')
print(f'R^2: {round(r2_score(y_test,lasso_pred),2)}')
print()
print('Elastic-Net CV Regression:')
print(f'RMSE: {root_mean_squared_error(y_test,elastic_pred).round(2)}')
print(f'R^2: {round(r2_score(y_test,elastic_pred),2)}')
```

Linear Regression:

RMSE: 1.52

R^2: 0.92

Polynomial Regression, Degree 3:

RMSE: 0.58

R^2: 0.99

Ridge CV Regression:

RMSE: 0.62

R^2: 0.99

Lasso CV Regression:

RMSE: 0.61

R^2: 0.99

Elastic-Net CV Regression:

RMSE: 0.61

R^2: 0.99

```
In [47]: model_predictions_check = pd.DataFrame({'y_test':y_test,
                                              'Linear Regression':test_prediction.round(2),
                                              '3rd-Polynomial':final_pr_pred.round(2),
                                              'Ridge Regression':ridge_cv_pred_values.round(2),
                                              'Lasso Regression':lasso_pred.round(2),
                                              'Elastic-Net Regression':elastic_pred.round(2)})
```

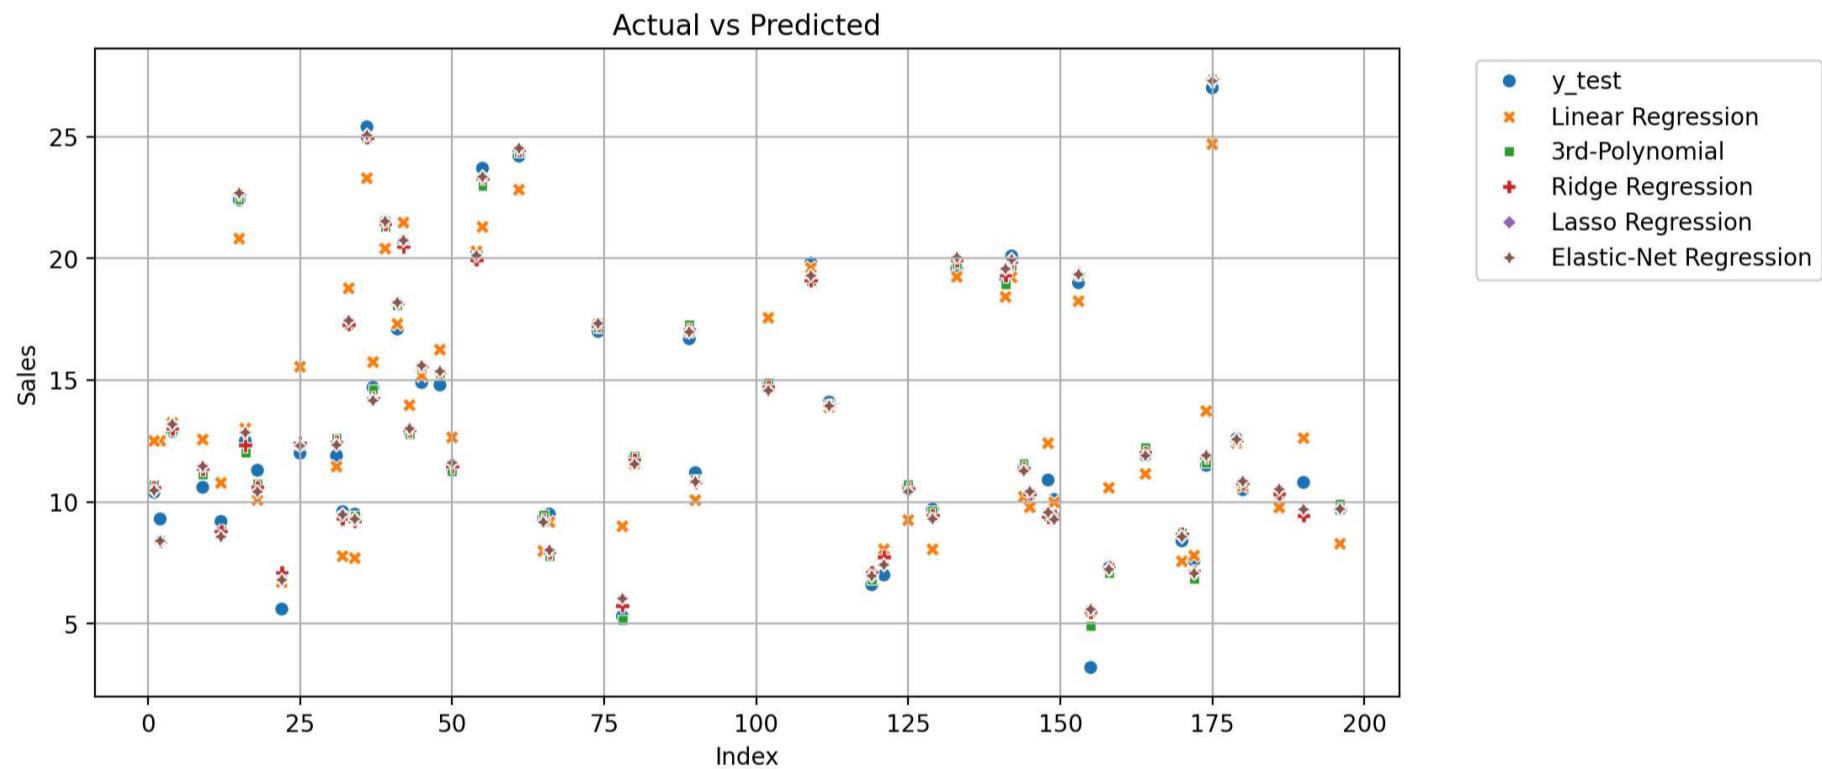
model_predictions_check

Out[47]:

| | y_test | Linear Regression | 3rd-Polynomial | Ridge Regression | Lasso Regression | Elastic-Net Regression |
|------------|--------|-------------------|----------------|------------------|------------------|------------------------|
| 37 | 14.7 | 15.74 | 14.62 | 14.25 | 14.16 | 14.16 |
| 109 | 19.8 | 19.61 | 19.06 | 19.10 | 19.29 | 19.29 |
| 31 | 11.9 | 11.45 | 12.61 | 12.45 | 12.34 | 12.34 |
| 89 | 16.7 | 17.01 | 17.29 | 17.07 | 16.98 | 16.98 |
| 66 | 9.5 | 9.17 | 7.79 | 7.94 | 8.02 | 8.02 |
| 119 | 6.6 | 7.01 | 6.82 | 7.10 | 6.96 | 6.96 |
| 54 | 20.2 | 20.29 | 19.96 | 19.95 | 20.12 | 20.12 |
| 74 | 17.0 | 17.30 | 17.13 | 17.25 | 17.33 | 17.33 |
| 145 | 10.3 | 9.78 | 10.40 | 10.34 | 10.43 | 10.43 |
| 142 | 20.1 | 19.22 | 19.71 | 19.86 | 19.96 | 19.96 |
| 148 | 10.9 | 12.41 | 9.41 | 9.38 | 9.57 | 9.57 |
| 112 | 14.1 | 13.89 | 13.99 | 13.96 | 13.94 | 13.94 |
| 174 | 11.5 | 13.73 | 11.62 | 11.85 | 11.91 | 11.91 |
| 55 | 23.7 | 21.29 | 22.98 | 23.29 | 23.35 | 23.35 |
| 141 | 19.2 | 18.42 | 18.95 | 19.30 | 19.57 | 19.57 |
| 149 | 10.1 | 9.98 | 9.41 | 9.42 | 9.28 | 9.28 |
| 25 | 12.0 | 15.55 | 12.33 | 12.42 | 12.30 | 12.30 |
| 34 | 9.5 | 7.69 | 9.43 | 9.22 | 9.30 | 9.30 |
| 170 | 8.4 | 7.56 | 8.75 | 8.70 | 8.57 | 8.57 |
| 39 | 21.5 | 20.40 | 21.28 | 21.42 | 21.52 | 21.52 |
| 172 | 7.6 | 7.79 | 6.83 | 7.15 | 7.07 | 7.07 |
| 153 | 19.0 | 18.24 | 19.30 | 19.36 | 19.34 | 19.34 |
| 175 | 27.0 | 24.69 | 27.39 | 27.33 | 27.28 | 27.28 |
| 61 | 24.2 | 22.82 | 24.37 | 24.45 | 24.52 | 24.52 |
| 65 | 9.3 | 7.98 | 9.48 | 9.24 | 9.17 | 9.17 |
| 50 | 11.4 | 12.65 | 11.25 | 11.46 | 11.61 | 11.61 |
| 42 | 20.7 | 21.47 | 20.73 | 20.47 | 20.74 | 20.74 |
| 129 | 9.7 | 8.05 | 9.64 | 9.46 | 9.30 | 9.30 |
| 179 | 12.6 | 12.42 | 12.52 | 12.55 | 12.57 | 12.57 |
| 2 | 9.3 | 12.51 | 8.43 | 8.39 | 8.40 | 8.40 |
| 12 | 9.2 | 10.78 | 8.82 | 8.79 | 8.57 | 8.57 |
| 133 | 19.6 | 19.24 | 19.76 | 19.92 | 20.04 | 20.04 |
| 90 | 11.2 | 10.07 | 10.92 | 10.81 | 10.83 | 10.83 |
| 22 | 5.6 | 6.71 | 6.97 | 7.11 | 6.80 | 6.80 |
| 41 | 17.1 | 17.31 | 18.08 | 18.16 | 18.19 | 18.19 |
| 32 | 9.6 | 7.77 | 9.45 | 9.31 | 9.48 | 9.48 |
| 125 | 10.6 | 9.25 | 10.72 | 10.54 | 10.42 | 10.42 |
| 196 | 9.7 | 8.28 | 9.94 | 9.71 | 9.71 | 9.71 |
| 158 | 7.3 | 10.58 | 7.08 | 7.31 | 7.23 | 7.23 |
| 180 | 10.5 | 10.64 | 10.77 | 10.77 | 10.85 | 10.85 |
| 16 | 12.5 | 13.01 | 12.03 | 12.33 | 12.85 | 12.85 |
| 186 | 10.3 | 9.77 | 10.34 | 10.34 | 10.52 | 10.52 |
| 144 | 11.4 | 10.21 | 11.59 | 11.39 | 11.27 | 11.27 |
| 121 | 7.0 | 8.05 | 7.65 | 7.75 | 7.42 | 7.42 |
| 80 | 11.8 | 11.57 | 11.89 | 11.73 | 11.55 | 11.55 |
| 18 | 11.3 | 10.08 | 10.75 | 10.61 | 10.42 | 10.42 |
| 78 | 5.3 | 9.00 | 5.15 | 5.74 | 6.03 | 6.03 |
| 48 | 14.8 | 16.25 | 15.28 | 15.33 | 15.36 | 15.36 |

| y_test | Linear Regression | 3rd-Polynomial | Ridge Regression | Lasso Regression | Elastic-Net Regression |
|--------|-------------------|----------------|------------------|------------------|------------------------|
| 4 | 12.9 | 13.24 | 12.96 | 13.01 | 13.19 |
| 15 | 22.4 | 20.81 | 22.43 | 22.65 | 22.68 |
| 1 | 10.4 | 12.50 | 10.72 | 10.60 | 10.47 |
| 43 | 12.9 | 13.97 | 12.78 | 12.95 | 13.01 |
| 102 | 14.8 | 17.56 | 14.90 | 14.72 | 14.57 |
| 164 | 11.9 | 11.15 | 12.25 | 12.02 | 11.90 |
| 9 | 10.6 | 12.56 | 11.13 | 11.36 | 11.47 |
| 155 | 3.2 | 5.51 | 4.92 | 5.48 | 5.59 |
| 36 | 25.4 | 23.29 | 24.95 | 24.96 | 25.06 |
| 190 | 10.8 | 12.62 | 9.50 | 9.43 | 9.69 |
| 33 | 17.4 | 18.77 | 17.29 | 17.28 | 17.45 |
| 45 | 14.9 | 15.19 | 15.54 | 15.59 | 15.60 |

```
In [48]: plt.figure(figsize=(10,5),dpi=200)
sns.scatterplot(data=model_predictions_check)
plt.title('Actual vs Predicted')
plt.xlabel('Index')
plt.ylabel('Sales')
plt.legend(bbox_to_anchor=(1.05,1))
plt.grid(True)
plt.show()
```



From the results, we can understand that Polynomial Regression performed better, and there wasnt overfitting since the results of Polyonal Regression and Regularizations are so close to each other.

```
In [ ]:
```